

Aim

1. Clean and Featurize data

```
In [1]: import html
import os
import re
import string
import unicodedata
from datetime import datetime

import category_encoders as ce
import numpy as np
import pandas as pd
from bs4 import BeautifulSoup
from fuzzywuzzy import fuzz
from nltk.corpus import stopwords
from pandarallel import pandarallel

# from gensim.corpora import Dictionary
# from gensim.models import TfidfModel
# from gensim.utils import simple_preprocess
# from sklearn.feature_extraction.text import TfidfVectorizer
from tqdm import tqdm
```

```
In [2]: DATA_ROOT = "../data"

# os.makedirs(f"{DATA_ROOT}/train/features")
```

```
In [3]: df_train = pd.read_pickle(f"{DATA_ROOT}/train/raw/data.pkl")
df_test = pd.read_pickle(f"{DATA_ROOT}/test/raw/data.pkl")
```

```
In [4]: df_train.head(2)
```

```
Out[4]:
```

	ID	Source	TMC	Start_Time	Distance(mi)	Description	Side	City	C
0	A-2478859	Bing	0.0	2016-02-08 00:37:08	3.23	Between Sawmill Rd/Exit 20 and OH-315/Olentang...	R	Dublin	F
1	A-1	MapQuest	201.0	2016-02-08 05:46:00	0.01	Right lane blocked due to accident on I-70 Eas...	R	Dayton	Montg

2 rows × 38 columns

```
In [5]: df_test.head(2)
```

Out[5]:	ID	Source	TMC	Start_Time	Distance(mi)	Description	Side	City	County
0	A-3017746	Bing	0.0	2020-01-01 00:01:00	0.0	At Hampshire Rd/Exit 41 - Accident.	R	Westlake Village	Ventura
1	A-3017745	Bing	0.0	2020-01-01 00:02:00	0.0	At Sheep Creek Rd - Accident.	L	Phelan	San Bernardino

2 rows x 38 columns

```
In [6]: df_train.columns
```

```
Out[6]: Index(['ID', 'Source', 'TMC', 'Start_Time', 'Distance(mi)', 'Description',
            'Side', 'City', 'County', 'State', 'Zipcode', 'Timezone',
            'Airport_Code', 'Temperature(F)', 'Humidity(%)', 'Pressure(in)',
            'Visibility(mi)', 'Wind_Direction', 'Wind_Speed(mph)',
            'Weather_Condition', 'Amenity', 'Bump', 'Crossing', 'Give_Way',
            'Junction', 'No_Exit', 'Railway', 'Roundabout', 'Station', 'Stop',
            'Traffic_Calming', 'Traffic_Signal', 'Turning_Loop', 'Sunrise_Sunset',
            'Civil_Twilight', 'Nautical_Twilight', 'Astronomical_Twilight',
            'Severity'],
            dtype='object')
```

Cleaning: Wind_Direction

```
In [7]: df_train["Wind_Direction"].unique()
```

```
Out[7]: array(['SW', 'Calm', 'SSW', 'WSW', 'WNW', 'NW', 'West', 'NNW', 'NNE',
            'South', 'W', 'North', 'Variable', 'SSE', 'SE', 'ESE', 'none',
            'East', 'NE', 'ENE', 'E', 'CALM', 'S', 'VAR', 'N'], dtype=object)
```

1. convert to upper case
2. replace VAR with variable

```
In [8]: def clean_wind_direction(df):
        df["Wind_Direction"] = df["Wind_Direction"].str.upper()
        df["Wind_Direction"] = df["Wind_Direction"].apply(
            lambda x: "variable" if x == "VAR" else x
        )

        return df

df_train = clean_wind_direction(df_train)
df_test = clean_wind_direction(df_test)
```

```
In [9]: print(sorted(df_train["Wind_Direction"].unique()))
```

```
['CALM', 'E', 'EAST', 'ENE', 'ESE', 'N', 'NE', 'NNE', 'NNW', 'NONE', 'NORTH',
 'NW', 'S', 'SE', 'SOUTH', 'SSE', 'SSW', 'SW', 'VARIABLE', 'W', 'WEST',
 'WNW', 'WSW', 'variable']
```

```
In [10]: print(sorted(df_test["Wind_Direction"].unique()))
```

```
['CALM', 'E', 'ENE', 'ESE', 'N', 'NE', 'NNE', 'NNW', 'NONE', 'NW', 'S', 'SE', 'SSE', 'SSW', 'SW', 'W', 'WNW', 'WSW', 'variable']
```

Cleaning: Description

```
In [18]: def rm_numbers(x):
          x = re.sub(r"[0-9]+", "", x)
          return x

def rm_html(x):
    x = html.unescape(x)
    x = BeautifulSoup(x).get_text()
    return x

def rm_url(x):
    x = re.sub("http*\s+", " ", x)
    return x

def rm_multiple_dots(x):
    x = re.sub(r"\.+", ". ", x)
    x = re.sub(r"\|+", ". ", x)
    return x

def rm_unicode(x):
    x = unicodedata.normalize("NFKD", x)
    return x

def rm_punctuation(x):
    x = re.sub("[%s]" % re.escape(string.punctuation.replace(".", "")), " ",
    return x

def rm_spaces(x):
    x = re.sub(" +", " ", x)
    return x

def rm_word(x, word):
    x = x.replace(word, "")
    return x

def clean_text(input_string):
    ss = input_string
    ss = rm_html(ss)
    ss = rm_url(ss)
    ss = rm_punctuation(ss)
    ss = rm_multiple_dots(ss)
    ss = rm_unicode(ss)
    ss = rm_spaces(ss)
    # ss = rm_numbers(ss)
    ss = rm_word(ss, "\n")
    ss = rm_word(ss, "\t")
    ss = ss.strip()

    return ss
```

```
In [20]: pandarallel.initialize(verbose=True,)
df_train["Description"] = df_train["Description"].parallel_apply(
    lambda x: clean_text(x)
)
```

```
In [26]: df_train["Description"].sample(5).values
```

```
Out[26]: array(['Accident on US 101 Oregon Coast Hwy near Cedar St.',
               'Accident on Bingle Rd at Houston Rosslyn Rd.',
               'Accident on Six Forks Rd at Lead Mine Rd.',
               'Right hand shoulder blocked due to accident on I 210 Eastbound before Exit 19 CA 2.',
               'Between VA 619 Exit 150 and VA 234 Exit 152 Accident.'],
              dtype=object)
```

```
In [24]: pandarallel.initialize(verbose=True,)
df_test["Description"] = df_test["Description"].parallel_apply(lambda x: clean_text(x))
```

```
In [27]: df_test["Description"].sample(5).values
```

```
Out[27]: array(['At Old Hiway Accident.',
               'Lane blocked due to accident on I 385 Northbound near Exit 34 Butler Rd.',
               'Accident on MN 36 Westbound at CR 35 Hadley Ave.',
               'At Southwood Plantation Rd Accident.',
               'Right lane blocked due to accident on Sam Houston Tlwy Eastbound at I 45 Gulf Fwy Exit 32.'],
              dtype=object)
```

Featurizing: Description

1. Get K keywords
2. create a binary vector of dimension K
3. if presence of word_x mark that dimension 1

Why this feature?

1. this will allow us to capture important words describing an accident
2. these words in turn might help in identifying severity

- [YAKE github](#)
- [Reference - Key word extractor](#)

YAKE is a lightweight, unsupervised automatic keyword extraction method that relies on statistical text features extracted from individual documents to identify the most relevant keywords in the text.

```
In [13]: import yake

kw_extractor = yake.KeywordExtractor()
text = ".".join(
    df_train["Description"].sample(n=100_000).tolist()
) # extracting from a sample as compute intensive process
language = "en"
max_ngram_size = 1
deduplication_threshold = 0.1
numOfKeywords = 1000
```

```

custom_kw_extractor = yake.KeywordExtractor(
    lan=language,
    n=max_ngram_size,
    dedupLim=deduplication_threshold,
    top=numOfKeywords,
    features=None,
    stopwords=stopwords.words("english"),
)
keywords = custom_kw_extractor.extract_keywords(text)

```

```

In [28]: # The lower the score, the more relevant the keyword is.
keywords

```

```

Out[28]: [('Accident', 6.169566081677733e-09),
 ('Northbound', 1.3069501235051463e-07),
 ('Hwy', 2.738028950102289e-06),
 ('ramp', 4.8401705088886265e-06),
 ('slow', 8.777147725501163e-06),
 ('Trl', 0.00014915096836162105),
 ('Mopac', 0.001457774067552258),
 ('Okeechobee', 0.0015743939528227214),
 ('Brookshire', 0.0016599371751449834),
 ('Huntington', 0.002962632417734576),
 ('NYS', 0.005783750908517111),
 ('Fuqua', 0.009709526119156922),
 ('Middlefield', 0.055070995136181314),
 ('JFK', 0.07611004069260154),
 ('Cedarhurst', 0.1031640040545067),
 ('57-56', 0.12423963791402771),
 ('PGBT.', 0.3325493080742511),
 ('Rhinecliff', 0.6395068253572433),
 ('Chavaneaux', 0.6552093622826601),
 ('Gibsonburg', 0.677368534171853)]

```

```

In [30]: pd.to_pickle(
    keywords, f"{DATA_ROOT}/train/keywords.pkl",
)

```

```

In [34]: # using top 15 words
keywords_list = [i[0] for i in keywords[:15]]

```

```

In [35]: keywords_list

```

```

Out[35]: ['Accident',
 'Northbound',
 'Hwy',
 'ramp',
 'slow',
 'Trl',
 'Mopac',
 'Okeechobee',
 'Brookshire',
 'Huntington',
 'NYS',
 'Fuqua',
 'Middlefield',
 'JFK',
 'Cedarhurst']

```

```

In [45]: fuzz.partial_ratio("hello world 2", "hello world") # demo of partial_ratio

```

```

Out[45]: 100

```

```
In [57]: def get_kw_vec(x, kw_list):
vec = [fuzz.partial_ratio(i.lower(), x.lower()) for i in kw_list]
vec = np.array(vec)
vec = np.where(vec > 60, 1, 0).tolist()
return vec
```

```
In [58]: pandarallel.initialize(verbose=True,)
df_train["kw_vec"] = df_train["Description"].parallel_apply(
    lambda x: get_kw_vec(x, keywords_list)
)
```

```
In [59]: pandarallel.initialize(verbose=True,)
df_test["kw_vec"] = df_test["Description"].parallel_apply(
    lambda x: get_kw_vec(x, keywords_list)
)
```

```
In [62]: df_train["kw_vec"].sample(5).head()
```

```
Out[62]: 2428374    [1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
209306      [1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0]
849060      [1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
70600       [1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
1845646     [1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
Name: kw_vec, dtype: object
```

Featurizing: Zipcode

Reference encoding zipcode

Why breakdown zip code?

- zip codes are hierarchical in nature
- using first N digits of zipcode will give us some understanding of the region
- using next N digits will give understanding of next sub regions
- we also reduce the number of unique zipcodes this way
- this will enable model to learn better

```
In [64]: df_train["zip_02"] = df_train["Zipcode"].str[:2]
df_test["zip_02"] = df_test["Zipcode"].str[:2]
```

```
In [65]: df_train["zip_25"] = df_train["Zipcode"].str[2:5]
df_test["zip_25"] = df_test["Zipcode"].str[2:5]
```

Zip codes are of varying length that could be one of the features

`zip_len` could be one of the features

```
In [66]: df_train["zip_len"] = df_train["Zipcode"].apply(len)
df_test["zip_len"] = df_test["Zipcode"].apply(len)
```

We can observe some compound zip codes like

1. 43068-3402
2. 93401-8325
3. 60607-3612

`is_compound` could be one of the boolean variables

```
In [67]: df_train["zip_is_compound"] = df_train["Zipcode"].apply(lambda x: "-" in x)
df_test["zip_is_compound"] = df_test["Zipcode"].apply(lambda x: "-" in x)
```

```
df_train = pd.read_pickle(f"{DATA_ROOT}/train/featurized/data.pkl") df_test = pd.read_pickle(f"{DATA_ROOT}/test/featurized/data.pkl")
```

Handling categorical features

Scikit learn's package `dtrees` does not handle categorical values

- we need to encode categories into some kind of encoding, in order to train the model.
- [reference](#)

```
In [4]: df_train.dtypes
```

```
Out[4]: ID                object
Source                object
TMC                  float64
Distance(mi)         float64
Side                 object
City                 object
County               object
State                object
Timezone             object
Airport_Code          object
Temperature(F)        float64
Humidity(%)           float64
Pressure(in)          float64
Visibility(mi)         float64
Wind_Direction         object
Wind_Speed(mph)        float64
Weather_Condition      object
Amenity                bool
Bump                   bool
Crossing               bool
Give_Way               bool
Junction               bool
No_Exit                bool
Railway                bool
Roundabout             bool
Station                bool
Stop                   bool
Traffic_Calming         bool
Traffic_Signal          bool
Turning_Loop            bool
Sunrise_Sunset          object
Civil_Twilight           object
Nautical_Twilight        object
Astronomical_Twilight     object
kw_vec                  object
zip_02                  object
zip_25                  object
zip_len                  int64
zip_is_compound          bool
Severity                 int64
dtype: object
```

Category transformer

```
In [5]: # making list of cateogrical features that need encoding
# features that are as dtype string will need encoding

# categorical_feature -> encoding method
categorical_features = {
    "Source": "base_2", # 3 unique values
    "Side": "base_2", # 2 unique values
    "City": "base_4", # 11895 unique values
    "County": "base_4", # 1713 unique values
    "State": "base_2", # 49 unique values
    "Timezone": "base_2", # 4 unique values
    "Airport_Code": "base_4", # 2001 unique values
    "Wind_Direction": "base_4", # 24 unique values (after cleaning)
    "Weather_Condition": "base_4", # 127 unique values
    "Sunrise_Sunset": "base_2", # 2 unique values
    "Civil_Twilight": "base_2", # 2 unique values
    "Nautical_Twilight": "base_2", # 2 unique value
    "Astronomical_Twilight": "base_2", # 2 unique value
    #
    # engineered features
    "zip_02": "ordinal",
    "zip_25": "ordinal",
}
```

```
In [10]: def category_transformer(
df_train: pd.DataFrame, df_test: pd.DataFrame, categorical_features: dict
):
    mapping_dict = dict()

    for i in categorical_features:
        method = categorical_features[i]

        # get values which will be encoded
        index_values = df_train[i].drop_duplicates().values.tolist()

        if "base" in method:
            print(f"encoding {i} with {method}")
            baseN = int(method.split("_")[1])
            enc = ce.binary.BaseNEncoder(base=2)
            # get values for train
            t_train = enc.fit_transform(df_train[i])
            df_train[i] = t_train.values.tolist()

            # get values for test
            t_test = enc.transform(df_test[i])
            df_test[i] = t_test.values.tolist()

        if method == "ordinal":
            print(f"encoding {i} as {method}")
            enc = ce.ordinal.OrdinalEncoder()
            # get values for train
            t_train = enc.fit_transform(df_train[i])
            df_train[i] = t_train.values.tolist()

            # get values for test
            t_test = enc.transform(df_test[i])
            df_test[i] = t_test.values.tolist()

    # store params
    mapping_dict[i] = {
        "enc_model_params": enc.get_params(),
```



```

        "enc_model_values": index_values,
    }

    return df_train, df_test, mapping_dict

df_train, df_test, mapping_dict = category_transformer(
    df_train, df_test, categorical_features
)

```

```

encoding Source with base_2
encoding Side with base_2
encoding City with base_4
encoding County with base_4
encoding State with base_2
encoding Timezone with base_2
encoding Airport_Code with base_4
encoding Wind_Direction with base_4
encoding Weather_Condition with base_4
encoding Sunrise_Sunset with base_2
encoding Civil_Twilight with base_2
encoding Nautical_Twilight with base_2
encoding Astronomical_Twilight with base_2
encoding zip_02 as ordinal
encoding zip_25 as ordinal

```

Mapping dictionary

```

In [13]: # sample of how categories are encoded
# -1 denote representation of unknown value
# -2 denotes representation of missing value
mapping_dict["Source"]

```

```

Out[13]: {'enc_model_params': {'base': 2,
    'cols': ['Source'],
    'drop_invariant': False,
    'handle_missing': 'value',
    'handle_unknown': 'value',
    'mapping': [{'col': 'Source',
        'mapping':
            Source_0  Source_1
            1         0         1
            2         1         0
            3         1         1
            -1        0         0
            -2        0         0}],
    'return_df': True,
    'verbose': 0},
    'enc_model_values': ['Bing', 'MapQuest', 'MapQuest-Bing']}

```

```

In [14]: # saving this as it will help in interpretation of results
pd.to_pickle(mapping_dict, f"{DATA_ROOT}/train/mapping_dict.pkl")

```

```

In [15]: boolean_features = {
    "Amenity",
    "Bump",
    "Crossing",
    "Give_Way",
    "Junction",
    "No_Exit",
    "Railway",
    "Roundabout",
    "Station",
    "Stop",

```

```

    "Traffic_Calming",
    "Traffic_Signal",
    "Turning_Loop",
    # engineered features
    "zip_is_compound",
}

```

```

In [ ]: len(
    set(df_train["Wind_Direction"].unique()).intersection(
        set(df_test["Wind_Direction"].unique())
    )
)

```

Out[]: 19

```

In [16]: final_feature_list = [
    "ID", # will be removing this before preparing modelling data
    "Source",
    "TMC",
    # "Start_Time", -> removing this as data is now sorted and split
    "Distance(mi)",
    # "Description", -> extracted features & removing this
    "Side",
    "City",
    "County",
    "State",
    # "Zipcode", -> extracted features & removing this
    "Timezone",
    "Airport_Code",
    "Temperature(F)",
    "Humidity(%)",
    "Pressure(in)",
    "Visibility(mi)",
    "Wind_Direction",
    "Wind_Speed(mph)",
    "Weather_Condition",
    "Amenity",
    "Bump",
    "Crossing",
    "Give_Way",
    "Junction",
    "No_Exit",
    "Railway",
    "Roundabout",
    "Station",
    "Stop",
    "Traffic_Calming",
    "Traffic_Signal",
    "Turning_Loop",
    "Sunrise_Sunset",
    "Civil_Twilight",
    "Nautical_Twilight",
    "Astronomical_Twilight",
    # engineered features
    "kw_vec",
    "zip_02",
    "zip_25",
    "zip_len",
    "zip_is_compound",
    # to predict
    "Severity",
]

```

```
In [17]: os.makedirs(f"{DATA_ROOT}/train/featurized/", exist_ok=True)
os.makedirs(f"{DATA_ROOT}/test/featurized/", exist_ok=True)
```

```
In [18]: df_train[final_feature_list].to_pickle(f"{DATA_ROOT}/train/featurized/data.p
df_test[final_feature_list].to_pickle(f"{DATA_ROOT}/test/featurized/data.pkl
```